

Building Agentic AI Applications with Databricks Agent Framework & MLflow 3

Technical Guide for the Multi-Country Pension Advisor

Repository: [pravinva/superannuation-agent-multi-country](#)

Architecture: ReAct Framework with Governance Layers

Deployment: Databricks Apps | Local Development

Table of Contents

1. [Architecture Overview](#)

2. [Core Components](#)

3. [Setup & Prerequisites](#)

4. [Building the Agent](#)

5. [Unity Catalog Integration](#)

6. [MLflow 3 Integration](#)

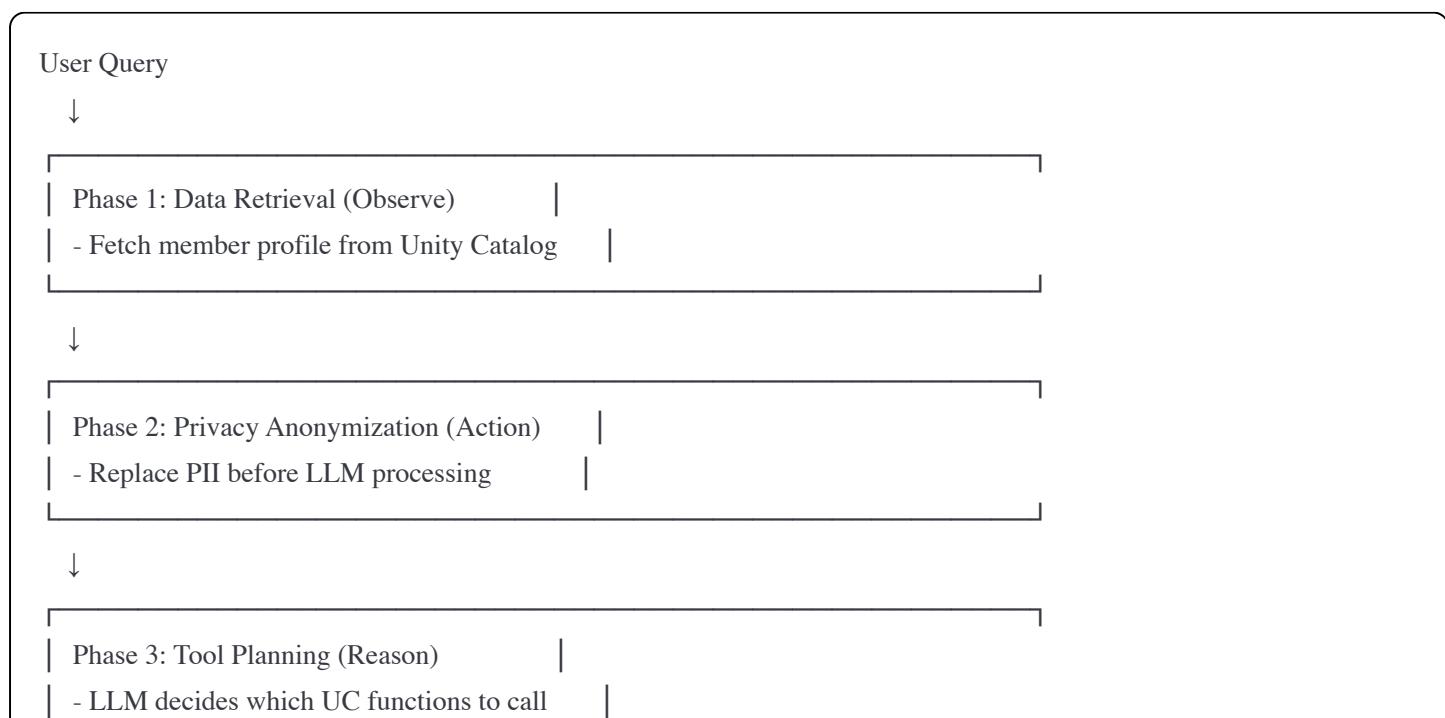
7. [Governance & Validation](#)

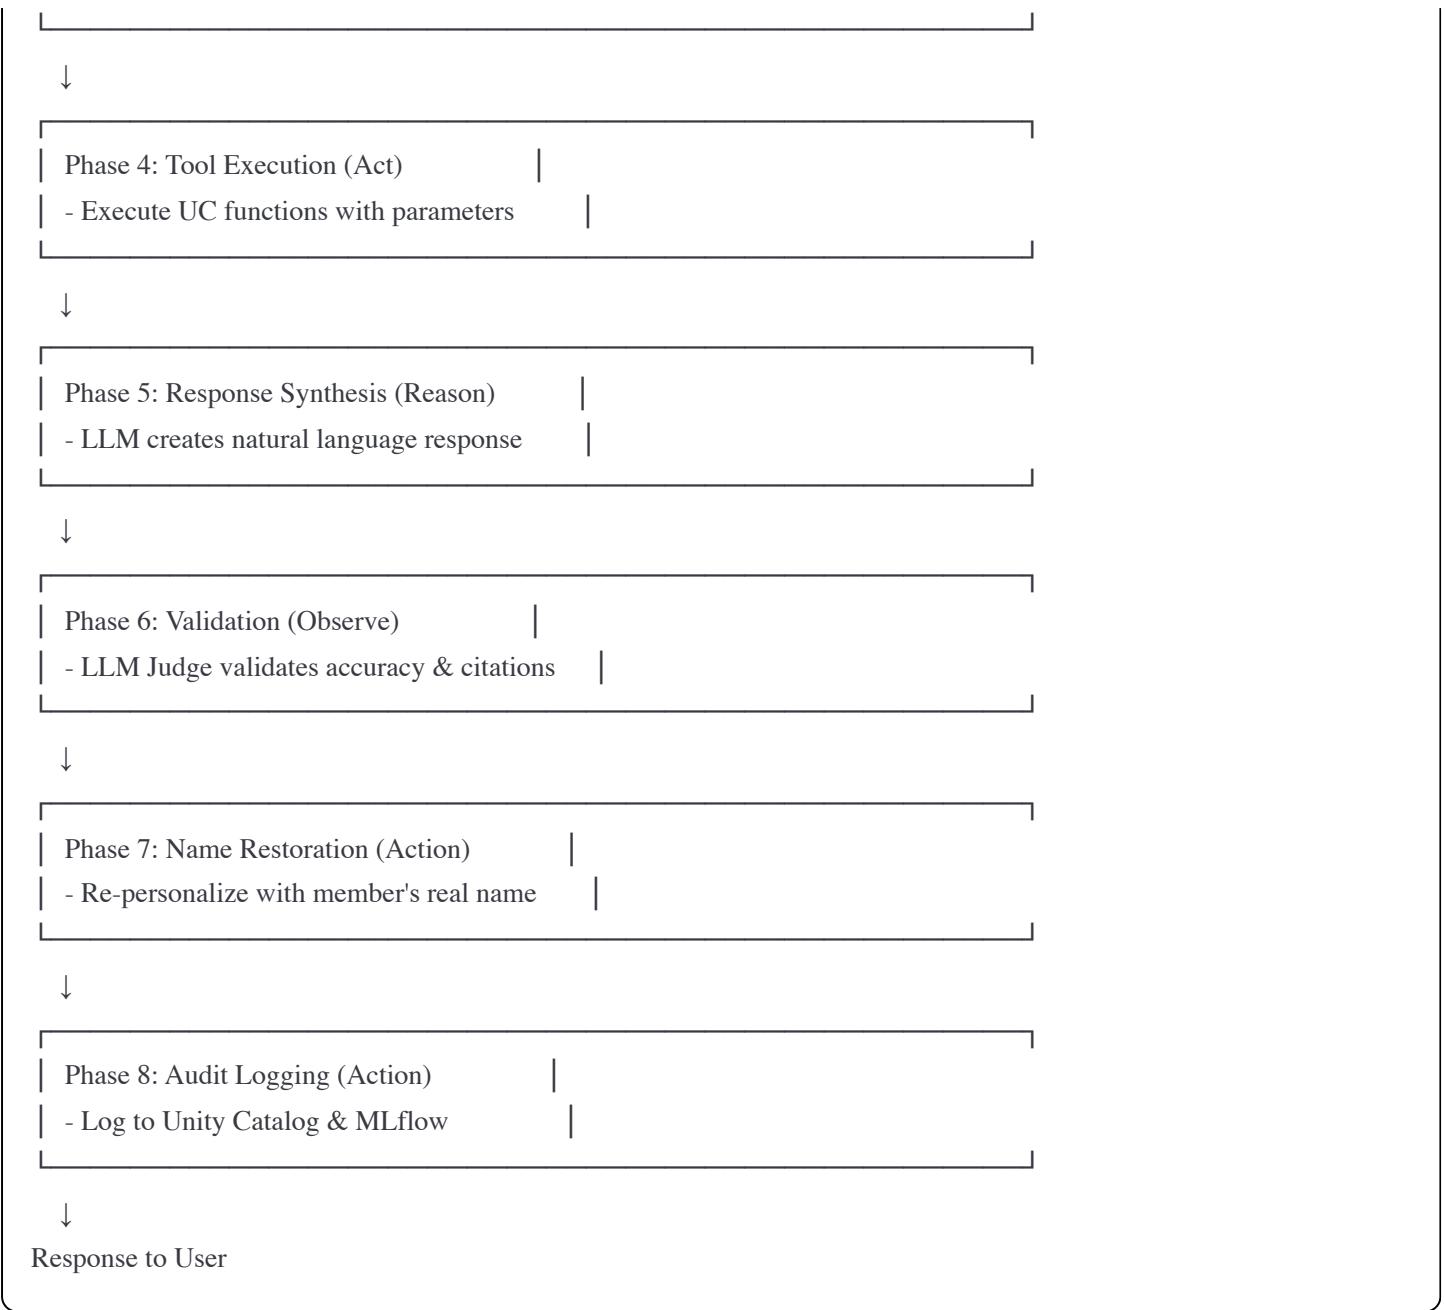
8. [Deployment Options](#)

9. [Code Examples](#)

Architecture Overview

The 8-Phase ReAct Pipeline





Key Design Principles

1. Separation of Concerns

- LLM handles reasoning and orchestration
- Unity Catalog functions handle calculations
- This prevents hallucinations in numerical results

2. ReAct Loop Implementation

- **Reason:** Plan which tools to use
- **Act:** Execute tools and observe results
- **Iterate:** Continue until task is complete

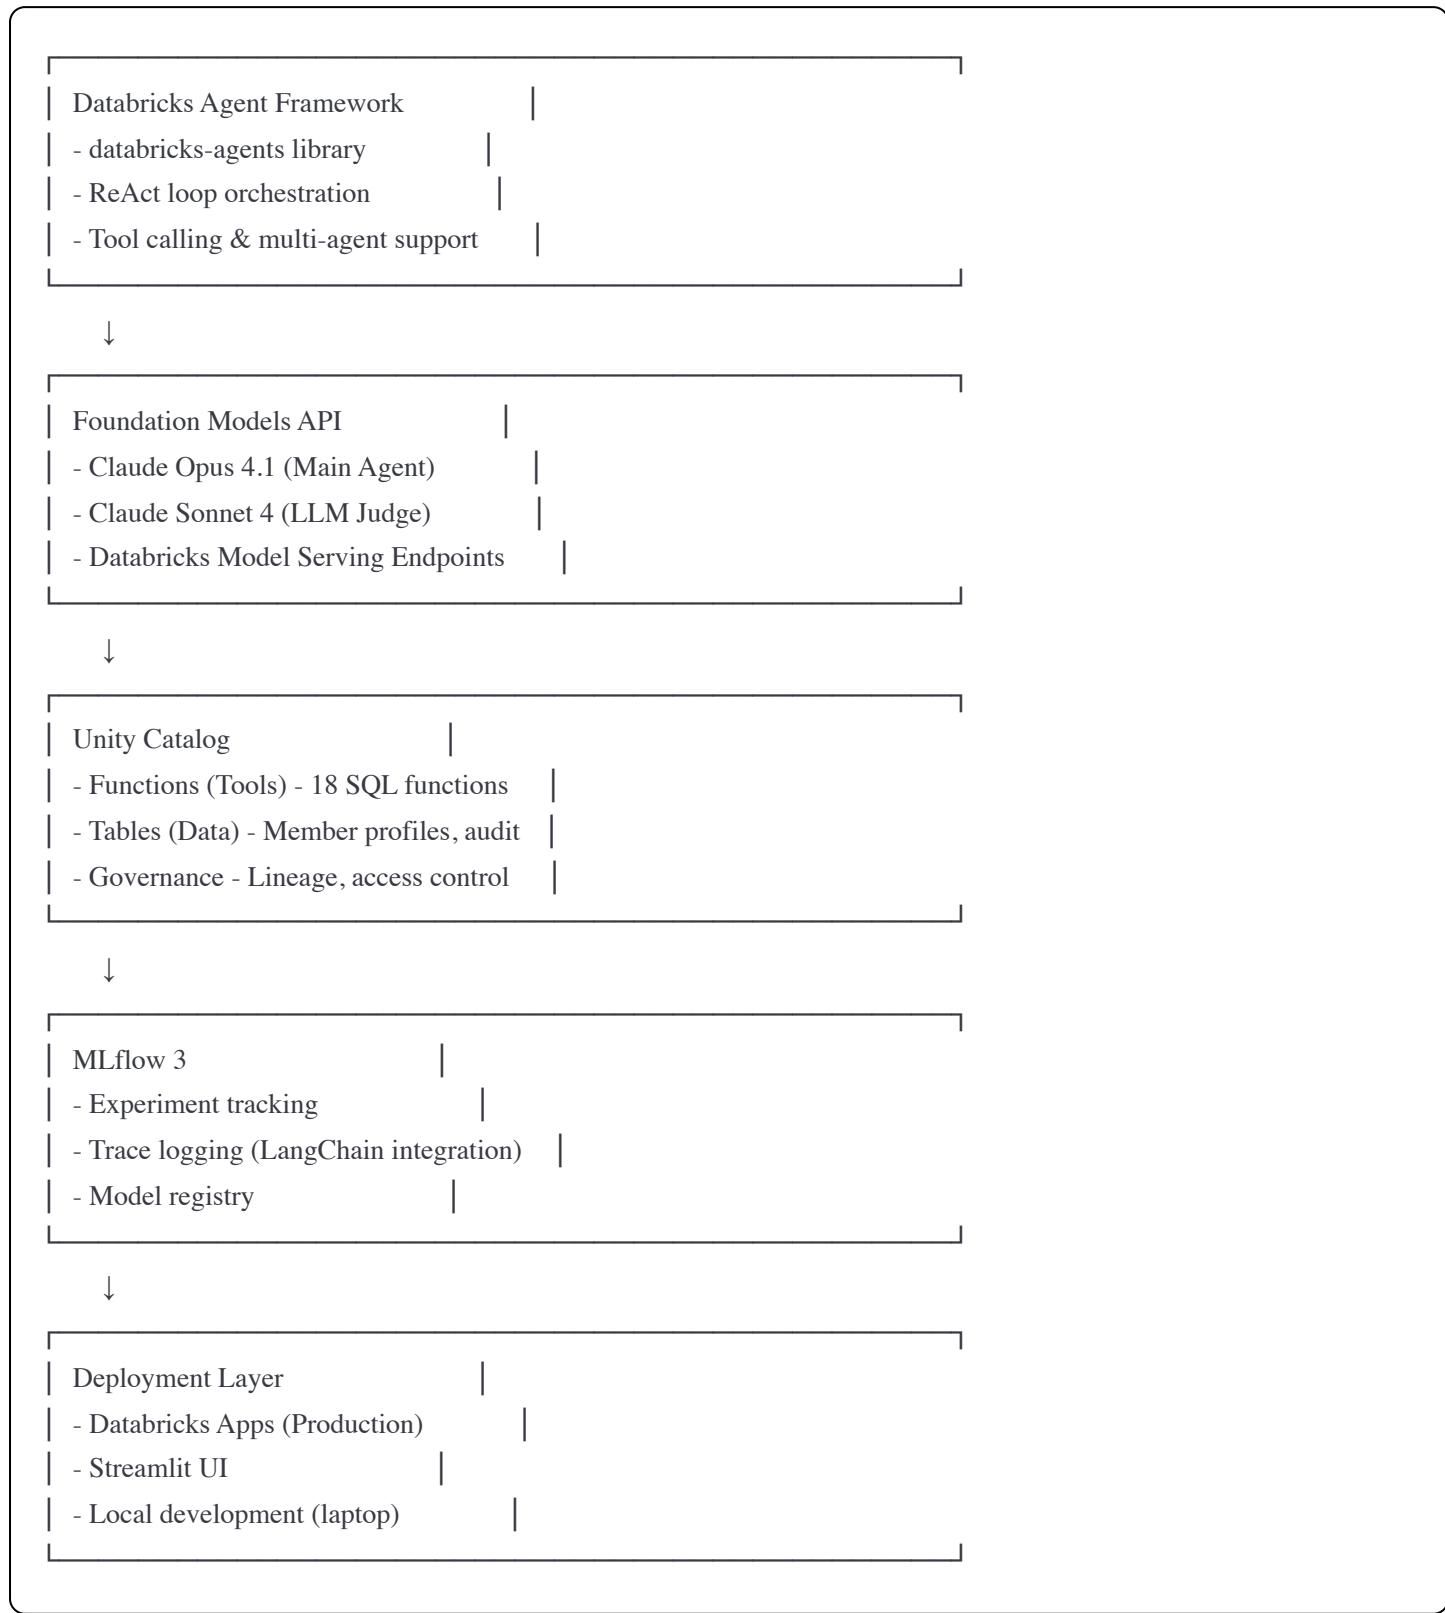
3. Governance by Design

- Privacy protection in Phase 2
- Quality validation in Phase 6

- Complete audit trail in Phase 8
-

Core Components

Technology Stack



Setup & Prerequisites

1. Environment Setup

Install Dependencies:

```
bash
```

```
# Core dependencies
pip install databricks-agents mlflow>=3.0.0 langchain-databricks

# UI dependencies (if building frontend)
pip install streamlit pandas

# Databricks SDK
pip install databricks-sdk
```

Project Structure:

```
superannuation-agent/
├── agent/
│   ├── __init__.py
│   ├── pension_agent.py      # Main agent implementation
│   ├── tools.py              # UC function wrappers
│   ├── privacy.py            # Anonymization logic
│   └── validation.py         # LLM Judge implementation
├── uc_functions/
│   ├── tax_calculators.sql   # Tax calculation functions
│   ├── benefit_checkers.sql  # Benefit eligibility
│   └── projections.sql       # Balance projections
├── config/
│   ├── agent_config.yaml     # Agent configuration
│   └── endpoints.yaml        # Model endpoint configs
├── ui/
│   └── app.py                # Streamlit interface
└── tests/
    └── test_agent.py
└── requirements.txt
└── databricks.yml           # Databricks App config
└── README.md
```

2. Unity Catalog Setup

Create Schema:

```
sql
```

```
CREATE SCHEMA IF NOT EXISTS main.pension_advisory;
```

-- Member profiles table

```
CREATE TABLE main.pension_advisory.members (
    member_id STRING,
    name STRING,
    age INT,
    balance DECIMAL(18,2),
    employment_status STRING,
    country STRING,
    years_of_service INT
);
```

-- Audit log table

```
CREATE TABLE main.pension_advisory.audit_log (
    interaction_id STRING,
    timestamp TIMESTAMP,
    member_id STRING,
    query STRING,
    response STRING,
    cost_usd DECIMAL(10,6),
    validation_verdict STRING,
    tools_called ARRAY<STRING>
);
```

3. Configure Model Serving Endpoints

Via Databricks UI:

1. Navigate to **Serving → Foundation Model API**
2. Enable Claude Opus 4.1 and Claude Sonnet 4
3. Note endpoint URLs (format: `/serving-endpoints/databricks-claude-opus-4`)

Via API:

```
python

from databricks.sdk import WorkspaceClient

w = WorkspaceClient()

# The endpoints are pre-configured for Foundation Models
main_endpoint = "databricks-claude-opus-4"
judge_endpoint = "databricks-claude-sonnet-4"
```

Building the Agent

Step 1: Define Unity Catalog Functions as Tools

Create UC Functions:

```
sql
-- Example: Australia tax calculator
CREATE OR REPLACE FUNCTION main.pension_advisory.au_calculate_tax(
    withdrawal_amount DECIMAL(18,2),
    age INT,
    balance DECIMAL(18,2)
)
RETURNS STRUCT<
    tax_amount DECIMAL(18,2),
    net_amount DECIMAL(18,2),
    tax_rate DECIMAL(5,2),
    explanation STRING
>
LANGUAGE SQL
RETURN
CASE
    WHEN age >= 60 THEN
        named_struct(
            'tax_amount', 0.0,
            'net_amount', withdrawal_amount,
            'tax_rate', 0.0,
            'explanation', 'Tax-free as age 60+'
        )
    WHEN age >= 55 AND age < 60 THEN
        named_struct(
            'tax_amount', withdrawal_amount * 0.15,
            'net_amount', withdrawal_amount * 0.85,
            'tax_rate', 15.0,
            'explanation', 'Taxed at 15% preservation age'
        )
    ELSE
        named_struct(
            'tax_amount', withdrawal_amount * 0.20,
            'net_amount', withdrawal_amount * 0.80,
            'tax_rate', 20.0,
            'explanation', 'Early access tax applies'
        )
END;
```

Wrap Functions in Python:

python

```
# agent/tools.py
from databricks.sdk import WorkspaceClient
from langchain.tools import tool
from typing import Dict, Any
```

```
w = WorkspaceClient()
```

```
@tool
```

```
def calculate_au_tax(withdrawal_amount: float, age: int, balance: float) -> Dict[str, Any]:
```

```
    """

```

```
    Calculate Australian superannuation withdrawal tax.
```

Args:

withdrawal_amount: Amount to withdraw in AUD

age: Member's age

balance: Current super balance in AUD

Returns:

Dictionary with tax_amount, net_amount, tax_rate, explanation

```
    """

```

```
result = w.statement_execution.execute_statement(
```

```
    warehouse_id="your_warehouse_id",
```

```
    catalog="main",
```

```
    schema="pension_advisory",
```

```
    statement=f"""

```

```
    SELECT

```

```
        main.pension_advisory.au_calculate_tax(

```

```
            {withdrawal_amount},

```

```
            {age},

```

```
            {balance}

```

```
        ) as result

```

```
    """

```

```
).result.data_array[0][0]
```

```
return {

```

```
    "tax_amount": result.tax_amount,

```

```
    "net_amount": result.net_amount,

```

```
    "tax_rate": result.tax_rate,

```

```
    "explanation": result.explanation

```

```
}
```

```
# Create tool list
```

```
pension_tools = [
```

```
    calculate_au_tax,
```

```
    calculate_us_401k_tax,
```

```
    calculate_uk_pension_tax,
```

```
# ... 15 more functions
```

```
]
```

Step 2: Implement Privacy Anonymization

```
python
```

```
# agent/privacy.py
import re
from typing import Tuple

def anonymize_member_data(member_data: dict, query: str) -> Tuple[dict, str, str]:
    """
    Replace member name with placeholder before LLM processing.

```

Args:

member_data: Dictionary with member profile
query: User's query string

Returns:

Tuple of (anonymized_data, anonymized_query, original_name)

```
    """
    original_name = member_data.get("name", "Member")
    placeholder = "MEMBER_NAME_PLACEHOLDER"
```

```
# Anonymize data
anonymized_data = member_data.copy()
anonymized_data["name"] = placeholder
```

```
# Anonymize query if name appears
```

```
anonymized_query = re.sub(
    rf"\b{re.escape(original_name)}\b",
    placeholder,
    query,
    flags=re.IGNORECASE
)
```

```
return anonymized_data, anonymized_query, original_name
```

```
def restore_member_name(response: str, original_name: str) -> str:
    """

```

Restore member's real name in the response.

```
    return response.replace("MEMBER_NAME_PLACEHOLDER", original_name)
```

Step 3: Build the Agent with ReAct Loop

python

```

# agent/pension_agent.py
from langchain_databricks import ChatDatabricks
from langchain.agents import AgentExecutor, create_react_agent
from langchain.prompts import PromptTemplate
from typing import Dict, Any
import mlflow
from agent.tools import pension_tools
from agent.privacy import anonymize_member_data, restore_member_name

class PensionAgent:
    def __init__(
        self,
        main_endpoint: str = "databricks-claude-opus-4",
        judge_endpoint: str = "databricks-claude-sonnet-4",
        catalog: str = "main",
        schema: str = "pension_advisory"
    ):
        self.catalog = catalog
        self.schema = schema

        # Main LLM for reasoning
        self.llm = ChatDatabricks(
            endpoint=main_endpoint,
            temperature=0.1,
            max_tokens=2000
        )

        # Judge LLM for validation
        self.judge_llm = ChatDatabricks(
            endpoint=judge_endpoint,
            temperature=0
        )

    # Create agent with tools
    self.agent = self._create_agent()

    def _create_agent(self) -> AgentExecutor:
        """Create ReAct agent with Unity Catalog tools."""

        # ReAct prompt template
        prompt = PromptTemplate.from_template("""
You are a retirement planning advisor. Answer the following question using the available tools.

Tools available:
{tools}
        """)

```

Tools available:

{tools}

Tool Names: {tool_names}

Member Context:

{member_context}

Question: {input}

Important:

- Always cite specific regulations (e.g., ATO rules, IRS code)
- Include appropriate disclaimers
- Use tools for ALL calculations - do not calculate in your head
- Be precise with numbers from tool outputs

Thought: {agent_scratchpad}

""")

```
# Create ReAct agent
agent = create_react_agent(
    llm=self.llm,
    tools=pension_tools,
    prompt=prompt
)

return AgentExecutor(
    agent=agent,
    tools=pension_tools,
    verbose=True,
    max_iterations=5,
    handle_parsing_errors=True
)

def process_query(
    self,
    member_id: str,
    query: str,
    validate: bool = True
) -> Dict[str, Any]:
    """
```

Process a member query through the 8-phase pipeline.

Args:

member_id: Member's unique identifier
query: Member's question
validate: Whether to run LLM Judge validation

Returns:

Dictionary with response, cost, validation results

```

"""
# Start MLflow run for tracking
with mlflow.start_run(run_name=f"query_{member_id}"):
    # PHASE 1: Data Retrieval
    member_data = self._retrieve_member_data(member_id)
    mlflow.log_param("member_id", member_id)
    mlflow.log_param("member_age", member_data["age"])
    mlflow.log_param("country", member_data["country"])

    # PHASE 2: Privacy Anonymization
    anon_data, anon_query, original_name = anonymize_member_data(
        member_data, query
    )
    mlflow.log_param("pii_anonymized", True)

    # PHASES 3-5: ReAct Loop (Reason → Act → Synthesize)
    # The agent handles tool planning, execution, and synthesis
    mlflow.log_param("query", query)

    response_data = self.agent.invoke({
        "input": anon_query,
        "member_context": self._format_member_context(anon_data)
    })

    response = response_data["output"]
    mlflow.log_text(response, "raw_response.txt")

    # Log tool calls
    intermediate_steps = response_data.get("intermediate_steps", [])
    tools_called = [step[0].tool for step in intermediate_steps]
    mlflow.log_param("tools_called", ",".join(tools_called))

    # PHASE 6: Validation
    validation_result = None
    if validate:
        validation_result = self._validate_response(
            query=query,
            response=response,
            tool_outputs=intermediate_steps
        )
    mlflow.log_dict(validation_result, "validation_result.json")

    # PHASE 7: Name Restoration
    final_response = restore_member_name(response, original_name)
    mlflow.log_text(final_response, "final_response.txt")

```

```

# Calculate costs
cost = self._calculate_cost(response_data, validation_result)
mlflow.log_metric("total_cost_usd", cost["total"])
mlflow.log_metric("main_llm_cost", cost["main_llm"])
mlflow.log_metric("judge_llm_cost", cost.get("judge_llm", 0))

# PHASE 8: Audit Logging
self._log_to_audit_table(
    member_id=member_id,
    query=query,
    response=final_response,
    cost=cost,
    validation=validation_result,
    tools_called=tools_called
)

return {
    "response": final_response,
    "cost": cost,
    "validation": validation_result,
    "tools_called": tools_called,
    "mlflow_run_id": mlflow.active_run().info.run_id
}

def _retrieve_member_data(self, member_id: str) -> Dict[str, Any]:
    """Fetch member profile from Unity Catalog."""
    from databricks.sdk import WorkspaceClient

    w = WorkspaceClient()
    result = w.statement_execution.execute_statement(
        warehouse_id="your_warehouse_id",
        catalog=self.catalog,
        schema=self.schema,
        statement=f"""
            SELECT * FROM {self.catalog}.{self.schema}.members
            WHERE member_id = '{member_id}'
        """
    )
    return result.data_array[0]

    return {
        "member_id": result[0],
        "name": result[1],
        "age": result[2],
        "balance": float(result[3]),
        "employment_status": result[4],
        "country": result[5],
    }

```

```
"years_of_service": result[6]
}

def _format_member_context(self, member_data: Dict) -> str:
    """Format member data for LLM context."""
    return f"""

Name: {member_data['name']}
Age: {member_data['age']}
Current Balance: ${member_data['balance']:.2f}
Employment: {member_data['employment_status']}
Country: {member_data['country']}
Years of Service: {member_data['years_of_service']}
"""


```

```
def _validate_response(
```

```
    self,
    query: str,
    response: str,
    tool_outputs: list
) -> Dict[str, Any]:
"""


```

LLM Judge validation.

Checks:

1. Response accuracy vs. tool outputs
2. Presence of citations
3. Appropriate disclaimers
4. No hallucinations

```
"""

validation_prompt = f"""


```

You are a quality validator for retirement advice responses.

Original Query: {query}

Tool Outputs: {tool_outputs}

Response to Validate: {response}

Evaluate:

1. Accuracy: Do numbers in response match tool outputs?
2. Citations: Are regulations cited (e.g., ATO, IRS)?
3. Disclaimers: Is appropriate disclaimer included?
4. Hallucinations: Any made-up information?

Respond in JSON:

```
{{
    "is_valid": true/false,
```

```
"confidence": 0.0-1.0,  
"violations": [],  
"reasoning": "explanation"  
}  
"""
```

```
validation_response = self.judge_llm.invoke(validation_prompt)
```

```
import json  
return json.loads(validation_response.content)
```

```
def _calculate_cost(  
    self,  
    response_data: Dict,  
    validation_result: Dict = None  
) -> Dict[str, float]:  
    """
```

Calculate costs based on Databricks Foundation Model API pricing.

Claude Opus 4.1: \$15 per 1M input tokens, \$75 per 1M output

Claude Sonnet 4: \$3 per 1M input tokens, \$15 per 1M output

```
"""
```

```
# Simplified cost calculation  
# In production, extract actual token counts from response metadata
```

```
main_llm_cost = 0.003 # Approximate per query  
judge_llm_cost = 0.0005 if validation_result else 0
```

```
return {  
    "main_llm": main_llm_cost,  
    "judge_llm": judge_llm_cost,  
    "total": main_llm_cost + judge_llm_cost  
}
```

```
def _log_to_audit_table(  
    self,  
    member_id: str,  
    query: str,  
    response: str,  
    cost: Dict,  
    validation: Dict,  
    tools_called: list  
):  
    """Log interaction to Unity Catalog audit table."""  
    from databricks.sdk import WorkspaceClient  
    import uuid  
    from datetime import datetime
```

```
w = WorkspaceClient()

interaction_id = str(uuid.uuid4())
timestamp = datetime.utcnow().isoformat()

w.statement_execution.execute_statement(
    warehouse_id="your_warehouse_id",
    catalog=self.catalog,
    schema=self.schema,
    statement=f"""
        INSERT INTO {self.catalog}.{self.schema}.audit_log
        VALUES (
            '{interaction_id}',
            '{timestamp}',
            '{member_id}',
            '{query.replace("", "")}',
            '{response.replace("", "")}',
            {cost['total']},
            '{validation.get("is_valid", "unknown")}',
            array({''.join([f'{t}' for t in tools_called])})
        )
    """
)
```

MLflow 3 Integration

Experiment Tracking

```
python

# Set experiment
mlflow.set_experiment("/Users/your_username/pension_advisor_experiments")

# The process_query() method already logs:
# - Parameters (member_id, age, country, query)
# - Metrics (cost, validation confidence)
# - Artifacts (responses, validation results)
# - Tags (tools_called)
```

Trace Logging with LangChain

MLflow 3 has native LangChain integration:

```
python
```

```
# Enable autologging
mlflow.langchain.autolog()

# All agent interactions are automatically traced:
# - LLM calls with prompts and responses
# - Tool invocations with inputs/outputs
# - Agent reasoning steps
# - Total duration and cost
```

Model Registry

```
python

# Register the agent as a model
with mlflow.start_run():

    mlflow.langchain.log_model(
        lc_model=agent,
        artifact_path="pension_agent",
        registered_model_name="pension_advisor_prod"
    )

# Load model for inference
loaded_agent = mlflow.langchain.load_model("models:/pension_advisor_prod/latest")
```

Governance & Validation

LLM Judge Implementation

The validation layer (Phase 6) uses a separate LLM to independently verify response quality:

```
python
```

```
# agent/validation.py
from typing import Dict, List, Any
from langchain_databricks import ChatDatabricks
import json

class LLMJudge:
    def __init__(self, endpoint: str = "databricks-claude-sonnet-4"):
        self.llm = ChatDatabricks(endpoint=endpoint, temperature=0)

    def validate(
        self,
        query: str,
        response: str,
        tool_outputs: List[Any],
        citations_required: bool = True
    ) -> Dict[str, Any]:
        """
        Validate response quality.

        Returns:
        {
            "is_valid": bool,
            "confidence": float,
            "violations": List[str],
            "reasoning": str
        }
        """
        prompt = self._create_validation_prompt(
            query, response, tool_outputs, citations_required
        )

        result = self.llm.invoke(prompt)

        try:
            validation = json.loads(result.content)
            return validation
        except json.JSONDecodeError:
            return {
                "is_valid": False,
                "confidence": 0.0,
                "violations": ["Failed to parse validation response"],
                "reasoning": result.content
            }

    def _create_validation_prompt(
```

```
self,
query: str,
response: str,
tool_outputs: List[Any],
citations_required: bool
) -> str:
    return f"""
```

You are an independent quality validator for financial advice responses.

Query: {query}

Tool Outputs:

```
{self._format_tool_outputs(tool_outputs)}
```

Response to Validate:

```
{response}
```

Validation Criteria:

1. ACCURACY: Numbers in response must match tool outputs exactly
2. CITATIONS: Must reference specific regulations (e.g., "ATO Section 307-5")
3. DISCLAIMERS: Must include appropriate risk disclaimers
4. HALLUCINATION: No information not derived from tools or known regulations
5. COMPLETENESS: Fully answers the query

Evaluate and respond in JSON format:

```
 {{
    "is_valid": true or false,
    "confidence": 0.0 to 1.0,
    "violations": ["list of any issues found"],
    "reasoning": "detailed explanation of validation decision"
}}
```

```
def _format_tool_outputs(self, tool_outputs: List[Any]) -> str:
    formatted = []
    for i, (action, observation) in enumerate(tool_outputs):
        formatted.append(f"Tool {i+1}: {action.tool}")
        formatted.append(f"Input: {action.tool_input}")
        formatted.append(f"Output: {observation}")
        formatted.append("---")
    return "\n".join(formatted)
```

Audit Trail Queries

Query the audit log for compliance:

```
sql
```

```
-- All interactions for a member
SELECT * FROM main.pension_advisory.audit_log
WHERE member_id = 'M12345'
ORDER BY timestamp DESC;
```

```
-- Failed validations
SELECT * FROM main.pension_advisory.audit_log
WHERE validation_verdict = 'false'
ORDER BY timestamp DESC;
```

```
-- High-cost queries
SELECT * FROM main.pension_advisory.audit_log
WHERE cost_usd > 0.01
ORDER BY cost_usd DESC;
```

```
-- Tool usage statistics
SELECT
    explode(tools_called) as tool_name,
    COUNT(*) as usage_count
FROM main.pension_advisory.audit_log
GROUP BY tool_name
ORDER BY usage_count DESC;
```

Deployment Options

Option 1: Databricks Apps (Production)

1. Create `databricks.yml`:

```
yaml
```

```

# databricks.yml

resources:
  apps:
    pension_advisor:
      name: pension-advisor-prod
      description: "Multi-Country Pension Advisor"

      # App configuration
      config:
        command:
          - "streamlit"
          - "run"
          - "ui/app.py"
          - "--server.port=8080"

      # Resources
      compute:
        size: MEDIUM
        auto_stop_mins: 30

      # Environment variables
      env:
        - name: DATABRICKS_HOST
          value: "{{workspace.host}}"
        - name: CATALOG
          value: "main"
        - name: SCHEMA
          value: "pension_advisory"

```

2. Deploy:

```

bash

# Authenticate
databricks auth login --host https://your-workspace.databricks.com

# Deploy app
databricks bundle deploy

# Start app
databricks bundle run pension_advisor

```

3. Access: App will be available at: <https://your-workspace.databricks.com/apps/pension-advisor-prod>

Option 2: Local Development (Laptop)

1. Set up authentication:

```
bash
```

```
# Option A: Use Databricks CLI  
databricks auth login --host https://your-workspace.databricks.com  
  
# Option B: Set environment variables  
export DATABRICKS_HOST="https://your-workspace.databricks.com"  
export DATABRICKS_TOKEN="your_personal_access_token"
```

2. Run locally:

```
bash
```

```
# Install dependencies  
pip install -r requirements.txt  
  
# Run Streamlit app  
streamlit run ui/app.py  
  
# Or run agent directly in Python  
python -c "  
from agent.pension_agent import PensionAgent  
  
agent = PensionAgent()  
result = agent.process_query(  
    member_id='M12345',  
    query='When can I access my pension?'  
)  
print(result['response'])  
"
```

3. Local development workflow:

```
python
```

```
# test_local.py
from agent.pension_agent import PensionAgent
import mlflow

# Set experiment for local testing
mlflow.set_experiment("/Users/your_name/pension_agent_local_dev")

# Initialize agent
agent = PensionAgent(
    main_endpoint="databricks-claude-opus-4",
    judge_endpoint="databricks-claude-sonnet-4"
)

# Test query
result = agent.process_query(
    member_id="TEST_001",
    query="What's my withdrawal tax?",
    validate=True
)

print(f"Response: {result['response']}")
print(f"Cost: ${result['cost']['total']:.4f}")
print(f"Valid: {result['validation']['is_valid']}")
print(f"MLflow Run: {result['mlflow_run_id']}")
```

Code Examples

Complete Streamlit UI

python

```
# ui/app.py
import streamlit as st
from agent.pension_agent import PensionAgent
import pandas as pd

st.set_page_config(
    page_title="Pension Advisor",
    page_icon="🏡",
    layout="wide"
)

# Initialize agent (cached)
@st.cache_resource
def get_agent():
    return PensionAgent()

agent = get_agent()

# Sidebar: Member selection
st.sidebar.header("Member Selection")

# Fetch members from UC
@st.cache_data
def load_members():
    from databricks.sdk import WorkspaceClient
    w = WorkspaceClient()
    result = w(statement_execution.execute_statement(
        warehouse_id="your_warehouse_id",
        catalog="main",
        schema="pension_advisory",
        statement="SELECT member_id, name, age, country FROM main.pension_advisory.members"
    ))
    return pd.DataFrame(
        result.result.data_array,
        columns=["member_id", "name", "age", "country"]
    )

members_df = load_members()
selected_member = st.sidebar.selectbox(
    "Select Member",
    members_df["member_id"],
    format_func=lambda x: members_df[members_df["member_id"] == x]["name"].values[0]
)

# Display member context
st.sidebar.subheader("Member Context")
```

```

member_info = members_df[members_df["member_id"] == selected_member].iloc[0]
st.sidebar.write(f"**Age:** {member_info['age']}")
st.sidebar.write(f"**Country:** {member_info['country']}")

# Main area: Query interface
st.title(" Retirement Planning Advisor")
st.markdown("Ask questions about your pension, tax implications, and withdrawal options.")

# Sample questions
st.subheader("Sample Questions")
col1, col2 = st.columns(2)
with col1:
    if st.button("When can I access my pension?"):
        query = "When can I access my pension?"
    if st.button("What are my tax implications?"):
        query = "What are my tax implications if I withdraw now?"
with col2:
    if st.button("What's my balance projection?"):
        query = "What will my balance be at retirement?"
    if st.button("Am I eligible for benefits?"):
        query = "Am I eligible for government benefits?"

# Text input
query = st.text_input("Or ask your own question:", "")

# Process query
if query:
    with st.spinner("Processing your query..."):

# Show pipeline progress
progress_bar = st.progress(0)
status_text = st.empty()

# Phases (for UI display)
phases = [
    "Retrieving member data...",
    "Protecting your privacy...",
    "Planning response...",
    "Calculating...",
    "Synthesizing answer...",
    "Validating accuracy...",
    "Finalizing...",
    "Logging for audit..."
]

for i, phase in enumerate(phases):
    status_text.text(phase)

```

```
progress_bar.progress((i + 1) / len(phases))

# Execute query
result = agent.process_query(
    member_id=selected_member,
    query=query,
    validate=True
)

progress_bar.progress(1.0)
status_text.text("Complete!")

# Display response
st.subheader("Response")
st.markdown(result["response"])

# Validation verdict
validation = result["validation"]
if validation["is_valid"]:
    st.success(f"✓ Response validated (Confidence: {validation['confidence']:.0%})")
else:
    st.warning(f"⚠ Validation issues: {', '.join(validation['violations'])}")

# Expandable: Details
with st.expander("View Details"):
    col1, col2, col3 = st.columns(3)
    with col1:
        st.metric("Cost", f"${result['cost']['total']:.4f}")
    with col2:
        st.metric("Tools Used", len(result["tools_called"]))
    with col3:
        st.metric("MLflow Run", result["mlflow_run_id"][:8] + "...")

    st.subheader("Tools Called")
    for tool in result["tools_called"]:
        st.write(f"- {tool}")

    st.subheader("Validation Details")
    st.json(validation)

# Governance tab
with st.expander("📊 Governance & Audit"):
    st.subheader("Recent Interactions")

# Query audit log
@st.cache_data
def load_audit_log(member_id, limit=10):
```

```
from databricks.sdk import WorkspaceClient
w = WorkspaceClient()
result = w.statement_execution.execute_statement(
    warehouse_id="your_warehouse_id",
    catalog="main",
    schema="pension_advisory",
    statement=f"""
        SELECT timestamp, query, cost_usd, validation_verdict
        FROM main.pension_advisory.audit_log
        WHERE member_id = '{member_id}'
        ORDER BY timestamp DESC
        LIMIT {limit}
    """
)
return pd.DataFrame(
    result.result.data_array,
    columns=["Timestamp", "Query", "Cost", "Valid"]
)

audit_df = load_audit_log(selected_member)
st.dataframe(audit_df, use_container_width=True)
```

Testing & Evaluation

Unit Tests

python

```
# tests/test_agent.py
import pytest
from agent.pension_agent import PensionAgent

@pytest.fixture
def agent():
    return PensionAgent()

def test_retrieve_member_data(agent):
    """Test member data retrieval."""
    member_data = agent._retrieve_member_data("M12345")
    assert "name" in member_data
    assert "age" in member_data
    assert "balance" in member_data

def test_privacy_anonymization():
    """Test PII anonymization."""
    from agent.privacy import anonymize_member_data, restore_member_name

    member_data = {"name": "John Smith", "age": 58}
    query = "When can John Smith retire?"

    anon_data, anon_query, original = anonymize_member_data(member_data, query)

    assert anon_data["name"] == "MEMBER_NAME_PLACEHOLDER"
    assert "John Smith" not in anon_query
    assert original == "John Smith"

    # Test restoration
    response = "MEMBER_NAME_PLACEHOLDER can retire at 60"
    restored = restore_member_name(response, original)
    assert "John Smith" in restored

def test_validation(agent):
    """Test LLM Judge validation."""
    result = agent._validate_response(
        query="What's my tax?",
        response="Your tax is 15% based on ATO rules.",
        tool_outputs=[("calculate_tax", "tax_rate: 15%")]
    )

    assert "is_valid" in result
    assert "confidence" in result
    assert "violations" in result

def test_end_to_end(agent):
```

```
"""Test full query processing."""
result = agent.process_query(
    member_id="TEST_001",
    query="When can I access my pension?",
    validate=True
)

assert "response" in result
assert "cost" in result
assert "validation" in result
assert result["cost"]["total"] > 0
```

Evaluation with MLflow

python

```
# evaluation/evaluate_agent.py
import mlflow
import pandas as pd
from agent.pension_agent import PensionAgent

# Load evaluation dataset
eval_data = pd.read_csv("evaluation/test_queries.csv")
# Columns: member_id, query, expected_answer, evaluation_criteria

agent = PensionAgent()

# Start evaluation run
with mlflow.start_run(run_name="agent_evaluation"):

    results = []

    for _, row in eval_data.iterrows():
        # Process query
        result = agent.process_query(
            member_id=row["member_id"],
            query=row["query"],
            validate=True
        )

        # Log individual result
        results.append({
            "query": row["query"],
            "response": result["response"],
            "cost": result["cost"]["total"],
            "is_valid": result["validation"]["is_valid"],
            "confidence": result["validation"]["confidence"]
        })

    # Calculate metrics
    results_df = pd.DataFrame(results)

    avg_cost = results_df["cost"].mean()
    validation_pass_rate = results_df["is_valid"].mean()
    avg_confidence = results_df["confidence"].mean()

    # Log metrics
    mlflow.log_metric("avg_cost_per_query", avg_cost)
    mlflow.log_metric("validation_pass_rate", validation_pass_rate)
    mlflow.log_metric("avg_validation_confidence", avg_confidence)

    # Log results artifact
```

```
mlflow.log_table(results_df, "evaluation_results.json")

print(f"Average Cost: ${avg_cost:.4f}")
print(f"Validation Pass Rate: {validation_pass_rate:.1%}")
print(f"Average Confidence: {avg_confidence:.1%}")
```

Best Practices

1. Cost Management

```
python
```

```
# Implement cost limits
class CostAwareAgent(PensionAgent):
    def __init__(self, max_cost_per_query: float = 0.01, *args, **kwargs):
        super().__init__(*args, **kwargs)
        self.max_cost_per_query = max_cost_per_query

    def process_query(self, *args, **kwargs):
        result = super().process_query(*args, **kwargs)

        if result["cost"]["total"] > self.max_cost_per_query:
            # Log alert
            mlflow.log_param("cost_exceeded", True)
            # Could trigger notification

        return result
```

2. Error Handling

```
python
```

```
# Graceful error handling
try:
    result = agent.process_query(member_id, query)
except Exception as e:
    # Log error
    mlflow.log_param("error", str(e))

    # Fallback response
    result = {
        "response": "I'm having trouble processing your query. Please contact support.",
        "error": str(e),
        "cost": {"total": 0}
    }
```

3. Caching for Performance

```
python
```

```
from functools import lru_cache

class CachedPensionAgent(PensionAgent):
    @lru_cache(maxsize=100)
    def _retrieve_member_data(self, member_id: str):
        """Cache member data for 1 hour."""
        return super().__retrieve_member_data(member_id)
```

4. A/B Testing

```
python
```

```
# Test different prompts or models
with mlflow.start_run(run_name="variant_a"):
    mlflow.log_param("variant", "detailed_prompt")
    agent_a = PensionAgent(prompt_version="detailed")
    result_a = agent_a.process_query(member_id, query)

with mlflow.start_run(run_name="variant_b"):
    mlflow.log_param("variant", "concise_prompt")
    agent_b = PensionAgent(prompt_version="concise")
    result_b = agent_b.process_query(member_id, query)

# Compare in MLflow UI
```

Troubleshooting

Common Issues

1. Authentication Errors

```
bash
```

```
# Verify authentication
databricks auth env
```

2. Unity Catalog Permissions

```
sql
```

```
-- Grant necessary permissions
GRANT USAGE ON CATALOG main TO `your_user`;
GRANT USAGE ON SCHEMA main.pension_advisory TO `your_user`;
GRANT SELECT ON TABLE main.pension_advisory.members TO `your_user`;
GRANT EXECUTE ON FUNCTION main.pension_advisory.au_calculate_tax TO `your_user`;
```

3. Model Serving Endpoint Not Found

```
python

# List available endpoints
from databricks.sdk import WorkspaceClient
w = WorkspaceClient()
endpoints = w.serving_endpoints.list()
for e in endpoints:
    print(f"{e.name}: {e.state.ready}")
```

4. MLflow Tracking URI

```
python

# Set tracking URI explicitly
mlflow.set_tracking_uri("databricks")
```

Performance Optimization

1. Batch Processing

```
python

def process_batch(member_ids: List[str], query: str):
    """Process same query for multiple members."""
    with mlflow.start_run(run_name="batch_processing"):
        results = []

        for member_id in member_ids:
            result = agent.process_query(member_id, query, validate=False)
            results.append(result)

        # Validate batch in parallel (if needed)
        # ...

    return results
```

2. Async Processing

```
python
```

```
import asyncio
from concurrent.futures import ThreadPoolExecutor

async def process_query_async(agent, member_id, query):
    """Async query processing."""
    loop = asyncio.get_event_loop()
    with ThreadPoolExecutor() as executor:
        result = await loop.run_in_executor(
            executor,
            agent.process_query,
            member_id,
            query
        )
    return result
```

Next Steps

1. Customize for Your Use Case

- Adapt UC functions for your regulatory environment
- Modify prompt templates
- Add domain-specific tools

2. Scale to Production

- Set up CI/CD with Databricks Asset Bundles
- Implement monitoring and alerting
- Configure auto-scaling

3. Enhance Capabilities

- Add more sophisticated validation logic
- Implement multi-agent collaboration
- Integrate with external data sources

4. Optimize Costs

- Implement caching strategies
- Use smaller models for simple queries
- Batch similar queries

Resources

- **Databricks Agent Framework:** <https://docs.databricks.com/en/generative-ai/agent-framework/index.html>
 - **MLflow 3 Documentation:** <https://mlflow.org/docs/latest/index.html>
 - **Unity Catalog Functions:** <https://docs.databricks.com/en/sql/language-manual/sql-ref-functions-udf-udaf.html>
 - **Databricks Apps:** <https://docs.databricks.com/en/dev-tools/databricks-apps/index.html>
 - **Foundation Model APIs:** <https://docs.databricks.com/en/machine-learning/foundation-models/index.html>
-

Summary

This architecture demonstrates:

- ✓ **ReAct Framework** - Reason → Act loop with governance
- ✓ **Separation of Concerns** - LLM orchestrates, UC functions calculate
- ✓ **Full Governance** - Privacy, validation, audit at every step
- ✓ **MLflow 3 Integration** - Complete observability and tracking
- ✓ **Production-Ready** - Deploy to Databricks Apps or run locally
- ✓ **Hallucination Mitigation** - Deterministic calculations prevent errors

The 8-phase pipeline ensures quality, compliance, and auditability while delivering instant, personalized retirement advice at scale.